

Measuring Memory Access Efficiency in Nsight Compute

Carl Pearson

github.com/cwpearson/nvidia-performance-tools
coalescing

Volta (GV100) Architecture

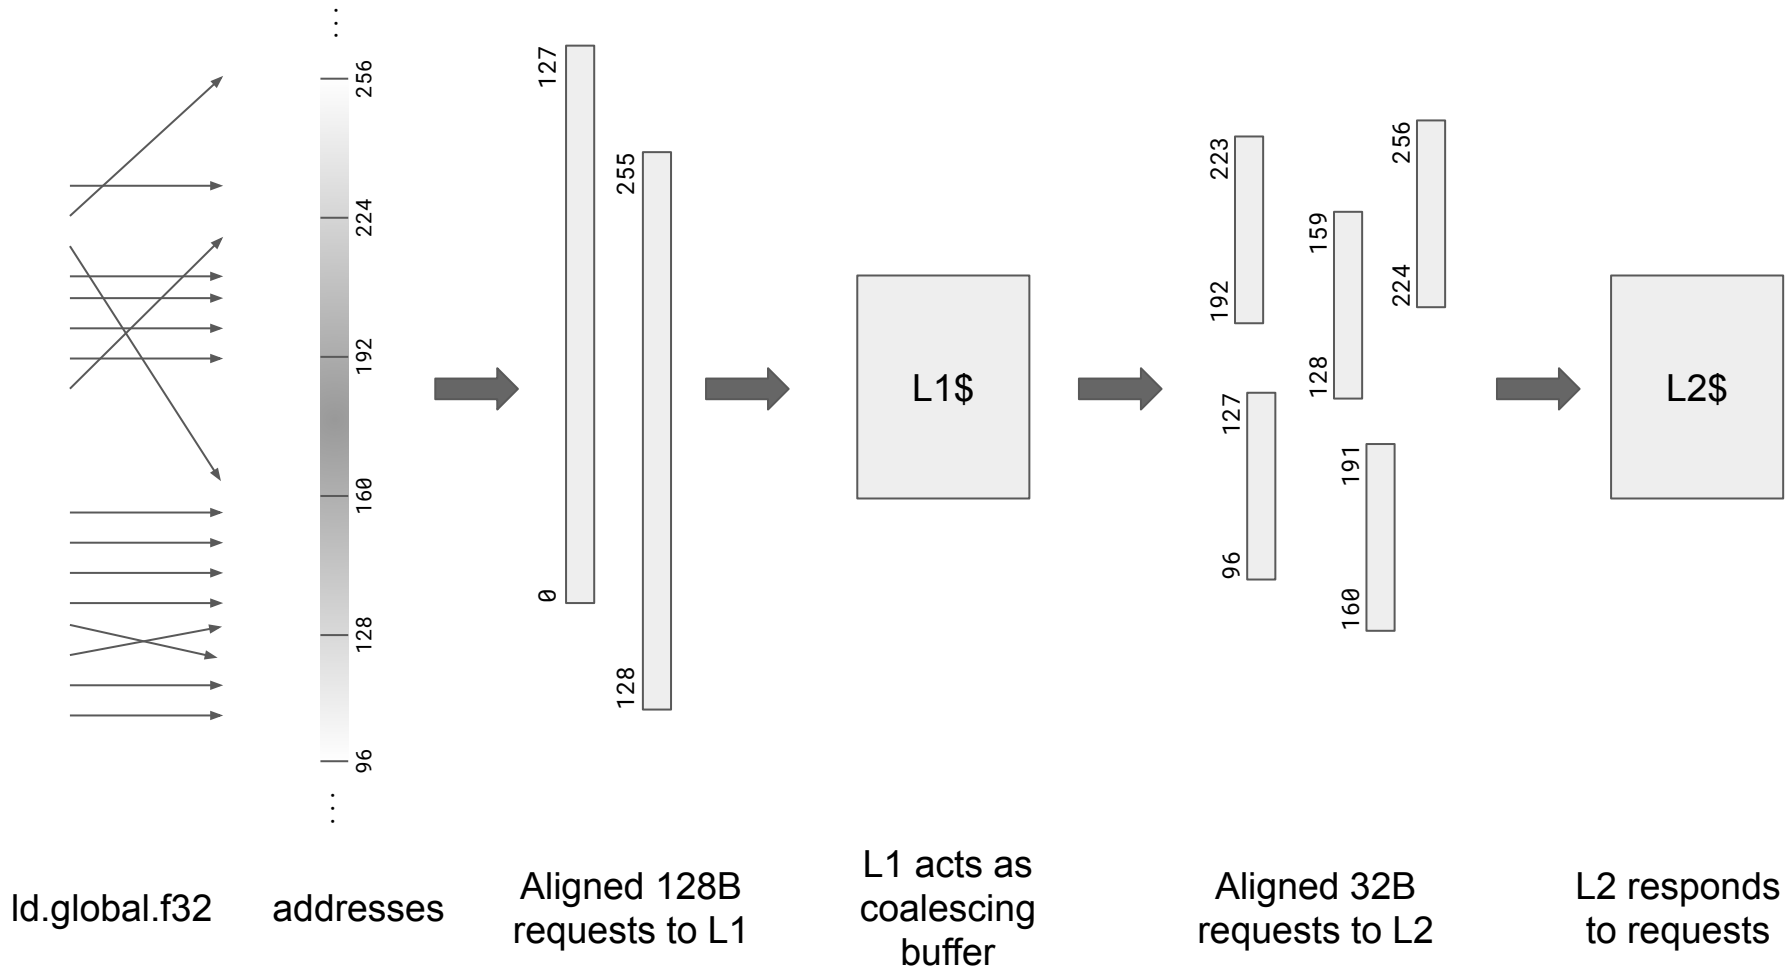
- Compute Capability 7.0
- Global memory is cached in L1^[4] and L2^[2]
- L1 is coalescing buffer^[1]
- Cache line is aligned 128B^[3]
- L1 accesses are 128B, L2 are 32B^[3]
- Warp memory requests are split into independent 128B requests^[3]

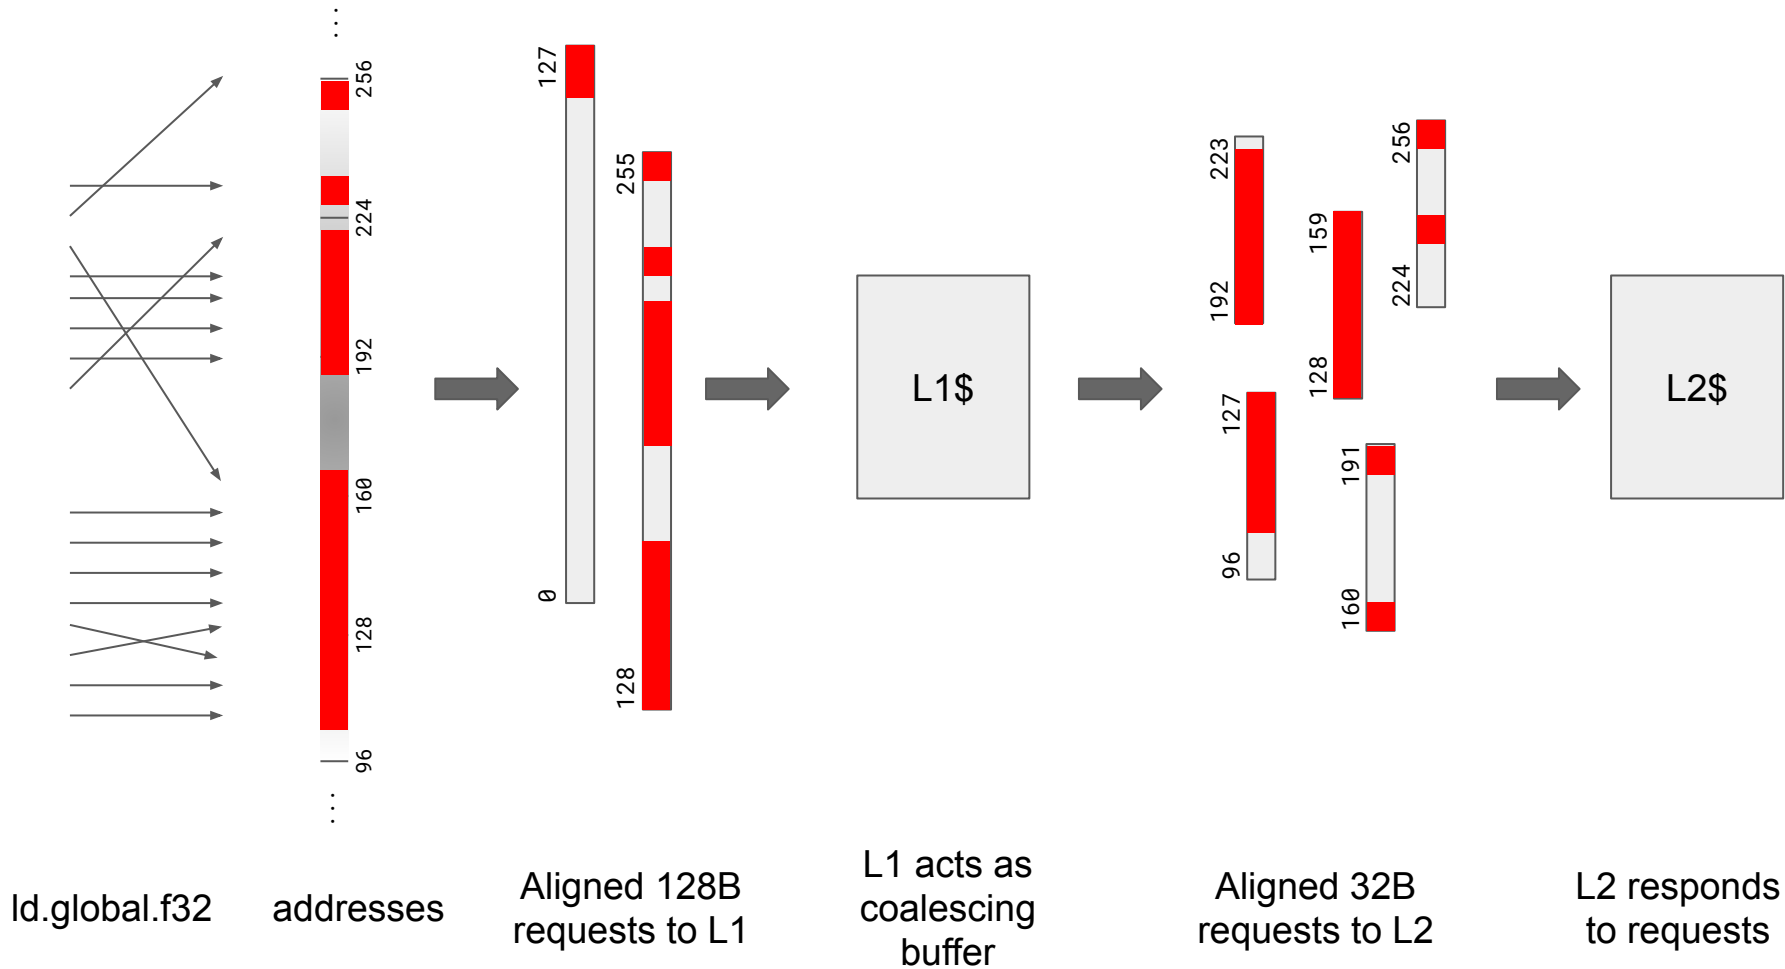
[1] [Volta Tuning Guide](#), 1.4.3.2 Unified Shared Memory/L1/Texture Cache

[2] [CUDA Programming Guide](#), H.4 Compute Capability 5.x


[3] [CUDA Programming Guide](#), H.3 Compute Capability 3.x

[4] [CUDA Programming Guide](#), H.6 Compute Capability 7.x






Setup (Coalesced)



10,000 = 10K offsets and elements (p).

Setup (Coalesced)



10,000 = 10K offsets and elements (p).

Source & PTX

```
int idx = off[i];  
float f = p[idx];  
f += 1;  
p[idx] = f;
```

```
int idx = off[i];  
double f = p[idx];  
f += 1;  
p[idx] = f;
```

```
ld.global.u32 %r9 [%rd10]  
ld.global.f32 %f1 [%rd12]  
add.f32 %f2, %f1, 0f3F800000  
std.global.f32 [%rd12] %f2
```

```
ld.global.u32 %r9 [%rd10]  
ld.global.f64 %fd1 [%rd12]  
add.f64 %fd2, %fd1, 0f3FF0000000000000  
std.global.f64 [%rd12] %fd2
```

Coalesced: (Titan V, GV100, CC7.0)

Source	PTX	Instrs. Executed	Predicated-On Thread Instrs. Executed	Mem Access Size	Mem L2 Transactions Global	Mem L1 Transactions Global
<code>float f = p[idx];</code>	<code>ld.global.f32 %f1 [%rd12]</code>	313	10,000	32	1250	313
<code>double f = p[idx];</code>	<code>ld.global.f64 %fd1 [%rd12]</code>	313	10,000	64	2500	625

10K instructions / (32 threads / warp) = 312.5 warps = 313 load instructions (312 full, 1 partial)

Each float load: 32 consecutive, aligned 4B regions

= 1 128B L1 request -> $313 * 1 = 313$ L1 transactions

= 4 32B L2 request -> $(312 * 4) + (1 * 2) = 1250$ L2 transactions

Each double load: 32 consecutive, aligned 8B regions

= 2 128B L1 request

= 8 32B L2 request

10K doubles = 80KB contiguous region

= $80\text{KB} / 128\text{B} = 625$ 128B L1 transactions = 2500 32B L2 transactions

Uncoalesced: (Titan V, GV100, CC7.0)

Source	PTX	Instrs. Executed	Predicated-On Thread Instrs. Executed	Mem Access Size	Mem L2 Transactions Global	Mem L1 Transactions Global
float f = p[idx];	ld.global.f32 %f1 [%rd12]	313	10,000	32	9895	9566
double f = p[idx];	ld.global.f64 %fd1 [%rd12]	313	10,000	64	9947	9740

Each float load: 32 random 4B regions

= 32 128B L1 request -> $313 * 32 = \sim 10\text{K}$ L1 transactions (some will randomly coalesce)

= 32 32B L2 request -> $313 * 31 = \sim 10\text{K}$ L2 transactions

Each double load: 32 random 8B regions

= 32 128B L1 request -> $313 * 32 = \sim 10\text{K}$ L1 transactions (some will randomly coalesce)

= 32 32B L2 request -> $313 * 31 = \sim 10\text{K}$ L2 transactions

Summary: Computing Memory Access Efficiency

Ideal L1 transactions per instruction

$$T_i = E * (S/8) / 128$$

E: (Predicated-On Thread Instructions Executed)

S: Memory Access Size (bits)


128: L1 transaction size (bytes)

T_e = Actual L1 transactions: measured by nsight

Efficiency: T_i / T_e

How to Reduce Efficiency

Unaligned:
2 requests:
128-255
256-383



Non-sequential:
3 requests:
0-127
128-255
256-383

